

# Data Transformation

```
library(nycflights13)
```

```
## Warning: package 'nycflights13' was built under R version 3.5.3
```

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 3.5.3
```

```
## -- Attaching packages -----  
----- tidyverse 1.3.0 --
```

```
## v ggplot2 3.3.2      v purrr   0.3.4  
## v tibble  3.0.3      v dplyr   1.0.0  
## v tidyr   1.1.0      v stringr 1.4.0  
## v readr   1.3.1      v forcats 0.5.0
```

```
## Warning: package 'readr' was built under R version 3.5.3
```

```
## Warning: package 'purrr' was built under R version 3.5.3
```

```
## Warning: package 'stringr' was built under R version 3.5.3
```

```
## Warning: package 'forcats' was built under R version 3.5.3
```

```
## -- Conflicts -----  
----- tidyverse_conflicts() --  
## x dplyr::filter() masks stats::filter()  
## x dplyr::lag()     masks stats::lag()
```

The data set “flights” will be used for the following examples. To learn more about the data set, please use `?flights`. The data set is in **tibble** format. The common variable types include -

- `int` stands for integers.
- `dbl` stands for doubles, or real numbers.
- `chr` stands for character vectors, or strings.
- `dtm` stands for date-times (a date + a time).
- `lgl` stands for logical, vectors that contain only TRUE or FALSE.

- `fctr` stands for factors, which R uses to represent categorical variables with fixed possible values.
- `date` stands for dates.

```
head(flights)
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2     830             819
## 2  2013     1     1     533             529           4     850             830
## 3  2013     1     1     542             540           2     923             850
## 4  2013     1     1     544             545          -1    1004            1022
## 5  2013     1     1     554             600          -6     812             837
## 6  2013     1     1     554             558          -4     740             728
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

## major dplyr functions

- Pick observations by their values `filter()` .
- Reorder the rows `arrange()` .
- Pick variables by their names `select()` .
- Create new variables with functions of existing variables `mutate()` .
- Collapse many values down to a single summary `summarise()` .
- change the scope `group_by()`

## filter()

print out the result and also assign the result to a new variable `dec25`.

```
(dec25 <- filter(flights, month == 12, day == 25))
```

```
## # A tibble: 719 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013    12    25     456           500          -4     649           651
## 2  2013    12    25     524           515           9     805           814
## 3  2013    12    25     542           540           2     832           850
## 4  2013    12    25     546           550          -4    1022          1027
## 5  2013    12    25     556           600          -4     730           745
## 6  2013    12    25     557           600          -3     743           752
## 7  2013    12    25     557           600          -3     818           831
## 8  2013    12    25     559           600          -1     855           856
## 9  2013    12    25     559           600          -1     849           855
## 10 2013    12    25     600           600           0     850           846
## # ... with 709 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```


**Notes** == brings problem when dealing with floating point numbers. Please use `near()` instead

```
1/49 * 49 == 1
```

```
## [1] FALSE
```

```
near(1/49 * 49, 1)
```

```
## [1] TRUE
```



R Markdown image

- use `|` and `&`, don't use `||` and `&&`.
- use `%in%` instead of `|` to simplify
- `!(x & y)` is the same as `!x | !y`
- `!(x | y)` is the same as `!x & !y`

```
nov_dec <- filter(flights, month %in% c(11, 12))  
  
# filter(flights, month == 11 | month == 12)
```

**missing values** is contagious.

```
NA > 5
```

```
## [1] NA
```

```
#> [1] NA  
10 == NA
```

```
## [1] NA
```

```
#> [1] NA  
NA + 10
```

```
## [1] NA
```

```
#> [1] NA  
NA / 2
```

```
## [1] NA
```

```
#> [1] NA  
NA == NA
```

```
## [1] NA
```

```
#> [1] NA
```

to check if the value is NA, `is.na()`

`filter()` only includes rows where the condition is TRUE; it excludes both FALSE and NA values. If you want to preserve missing values, ask for them explicitly:

```
df <- tibble(x = c(1, NA, 3))  
filter(df, x > 1)
```

```
## # A tibble: 1 x 1
##       x
##   <dbl>
## 1     3
```

```
filter(df, is.na(x) | x > 1)
```

```
## # A tibble: 2 x 1
##       x
##   <dbl>
## 1    NA
## 2     3
```

## arrange()

When sorting columns, NA are sorted at the end.

```
arrange(flights, year, month, day)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1     517             515         2     830             819
## 2  2013     1     1     533             529         4     850             830
## 3  2013     1     1     542             540         2     923             850
## 4  2013     1     1     544             545        -1    1004            1022
## 5  2013     1     1     554             600        -6     812             837
## 6  2013     1     1     554             558        -4     740             728
## 7  2013     1     1     555             600        -5     913             854
## 8  2013     1     1     557             600        -3     709             723
## 9  2013     1     1     557             600        -3     838             846
## 10 2013     1     1     558             600        -2     753             745
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
arrange(flights, desc(dep_delay))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     9     641           900       1301    1242         1530
## 2  2013     6    15    1432          1935       1137    1607         2120
## 3  2013     1    10    1121          1635       1126    1239         1810
## 4  2013     9    20    1139          1845       1014    1457         2210
## 5  2013     7    22     845          1600       1005    1044         1815
## 6  2013     4    10    1100          1900        960    1342         2211
## 7  2013     3    17    2321           810        911     135         1020
## 8  2013     6    27     959          1900        899    1236         2226
## 9  2013     7    22    2257           759        898     121         1026
## 10 2013    12     5     756          1700        896    1058         2020
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

## select()

```
select(flights, year, month, day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

```
# select(flights, year:day)
```

```
select(flights, -(year:day))
```

```
## # A tibble: 336,776 x 16
##   dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier
##   <int>         <int>         <dbl>   <int>         <int>         <dbl> <chr>
## 1     517           515           2     830           819           11 UA
## 2     533           529           4     850           830           20 UA
## 3     542           540           2     923           850           33 AA
## 4     544           545          -1    1004          1022          -18 B6
## 5     554           600          -6     812           837           -25 DL
## 6     554           558          -4     740           728           12 UA
## 7     555           600          -5     913           854           19 B6
## 8     557           600          -3     709           723           -14 EV
## 9     557           600          -3     838           846            -8 B6
## 10    558           600          -2     753           745            8 AA
## # ... with 336,766 more rows, and 9 more variables: flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

In addition, you can use **selection helpers**. Some helpers select specific columns:

- `starts_with()` : Starts with a prefix.
- `ends_with()` : Ends with a suffix.
- `contains()` : Contains a literal string.
- `matches()` : Matches a regular expression.
- `num_range()` : Matches a numerical range like x01, x02, x03.

```
select(flights, matches('time'))
```

```
## # A tibble: 336,776 x 6
##   dep_time sched_dep_time arr_time sched_arr_time air_time time_hour
##   <int>         <int>         <int>         <int>         <dbl> <dtm>
## 1     517           515           830           819         227 2013-01-01 05:00:00
## 2     533           529           850           830         227 2013-01-01 05:00:00
## 3     542           540           923           850         160 2013-01-01 05:00:00
## 4     544           545          1004          1022         183 2013-01-01 05:00:00
## 5     554           600           812           837         116 2013-01-01 06:00:00
## 6     554           558           740           728         150 2013-01-01 05:00:00
## 7     555           600           913           854         158 2013-01-01 06:00:00
## 8     557           600           709           723          53 2013-01-01 06:00:00
## 9     557           600           838           846         140 2013-01-01 06:00:00
## 10    558           600           753           745         138 2013-01-01 06:00:00
## # ... with 336,766 more rows
```

**\*\*rename\* columns**

```
rename(flights, tail_num = tailnum)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1     517           515         2     830           819
## 2  2013     1     1     533           529         4     850           830
## 3  2013     1     1     542           540         2     923           850
## 4  2013     1     1     544           545        -1    1004          1022
## 5  2013     1     1     554           600        -6     812           837
## 6  2013     1     1     554           558        -4     740           728
## 7  2013     1     1     555           600        -5     913           854
## 8  2013     1     1     557           600        -3     709           723
## 9  2013     1     1     557           600        -3     838           846
## 10 2013     1     1     558           600        -2     753           745
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tail_num <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

**everything()** move columns to the beginning

```
select(flights, time_hour, air_time, everything())
```

```
## # A tibble: 336,776 x 19
##   time_hour          air_time year month   day dep_time sched_dep_time
##   <dtm>            <dbl> <int> <int> <int>   <int>         <int>
## 1 2013-01-01 05:00:00      227  2013     1     1     517           515
## 2 2013-01-01 05:00:00      227  2013     1     1     533           529
## 3 2013-01-01 05:00:00      160  2013     1     1     542           540
## 4 2013-01-01 05:00:00      183  2013     1     1     544           545
## 5 2013-01-01 06:00:00      116  2013     1     1     554           600
## 6 2013-01-01 05:00:00      150  2013     1     1     554           558
## 7 2013-01-01 06:00:00      158  2013     1     1     555           600
## 8 2013-01-01 06:00:00       53  2013     1     1     557           600
## 9 2013-01-01 06:00:00      140  2013     1     1     557           600
## 10 2013-01-01 06:00:00      138  2013     1     1     558           600
## # ... with 336,766 more rows, and 12 more variables: dep_delay <dbl>,
## #   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, distance <dbl>,
## #   hour <dbl>, minute <dbl>
```



## mutate()

```
flights_sml <- select(flights,
  year:day,
  ends_with("delay"),
  distance,
  air_time
)
mutate(flights_sml,
  gain = dep_delay - arr_delay,
  speed = distance / air_time * 60
)
```

```
## # A tibble: 336,776 x 9
##   year month   day dep_delay arr_delay distance air_time   gain speed
##   <int> <int> <int>   <dbl>   <dbl>   <dbl>   <dbl> <dbl> <dbl>
## 1  2013     1     1         2       11    1400     227    -9   370.
## 2  2013     1     1         4       20    1416     227   -16   374.
## 3  2013     1     1         2       33    1089     160   -31   408.
## 4  2013     1     1        -1      -18    1576     183    17   517.
## 5  2013     1     1        -6      -25     762     116    19   394.
## 6  2013     1     1        -4       12     719     150   -16   288.
## 7  2013     1     1        -5       19    1065     158   -24   404.
## 8  2013     1     1        -3      -14     229      53    11   259.
## 9  2013     1     1        -3       -8     944     140     5   405.
## 10 2013     1     1        -2        8     733     138   -10   319.
## # ... with 336,766 more rows
```

```
# to keep the new variables only
transmute(flights,
  gain = dep_delay - arr_delay,
  hours = air_time / 60,
  gain_per_hour = gain / hours
)
```

```
## # A tibble: 336,776 x 3
##   gain hours gain_per_hour
##   <dbl> <dbl>         <dbl>
## 1     -9 3.78          -2.38
## 2    -16 3.78          -4.23
## 3    -31 2.67         -11.6
## 4     17 3.05           5.57
## 5     19 1.93           9.83
## 6    -16 2.5           -6.4
## 7    -24 2.63          -9.11
## 8     11 0.883         12.5
## 9      5 2.33           2.14
## 10   -10 2.3           -4.35
## # ... with 336,766 more rows
```

available operators include

- `+`, `-`, `^`
- `sum()`, `mean()`
- `/%` (integer division), `%%` (remainder)
- `log()`, `log2()`, `log10()`

```
transmute(flights,
  dep_time,
  hour = dep_time %/% 100,
  minute = dep_time %% 100
)
```

```
## # A tibble: 336,776 x 3
##   dep_time hour minute
##   <int> <dbl> <dbl>
## 1     517     5     17
## 2     533     5     33
## 3     542     5     42
## 4     544     5     44
## 5     554     5     54
## 6     554     5     54
## 7     555     5     55
## 8     557     5     57
## 9     557     5     57
## 10    558     5     58
## # ... with 336,766 more rows
```

- `lead()`, `lag()`

```
(x <- 1:10)
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
#> [1] 1 2 3 4 5 6 7 8 9 10  
lag(x)
```

```
## [1] NA 1 2 3 4 5 6 7 8 9
```

```
#> [1] NA 1 2 3 4 5 6 7 8 9  
lead(x)
```

```
## [1] 2 3 4 5 6 7 8 9 10 NA
```

```
#> [1] 2 3 4 5 6 7 8 9 10 NA
```

- `cumsum()`, `cumprod()`, `cummin()`, `cummax()`; and `dplyr` provides `cummean()` for cumulative means.

```
x
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
#> [1] 1 2 3 4 5 6 7 8 9 10  
cumsum(x)
```

```
## [1] 1 3 6 10 15 21 28 36 45 55
```

```
#> [1] 1 3 6 10 15 21 28 36 45 55  
cummean(x)
```

```
## [1] 1.000000 1.000000 1.333333 1.750000 2.200000 2.666667 3.142857 3.625000  
## [9] 4.111111 4.600000
```

- `min_rank()`, `row_number()`, `dense_rank()`, `percent_rank()`, `cume_dist()`, `ntile()`

```
y <- c(1, 2, 2, NA, 3, 4)  
min_rank(y)
```

```
## [1] 1 2 2 NA 4 5
```

```
#> [1] 1 2 2 NA 4 5  
min_rank(desc(y))
```

```
## [1] 5 3 3 NA 2 1
```

```
#> [1] 5 3 3 NA 2 1  
row_number(y)
```

```
## [1] 1 2 3 NA 4 5
```

```
#> [1] 1 2 3 NA 4 5  
dense_rank(y)
```

```
## [1] 1 2 2 NA 3 4
```

```
#> [1] 1 2 2 NA 3 4  
percent_rank(y)
```

```
## [1] 0.00 0.25 0.25 NA 0.75 1.00
```

```
#> [1] 0.00 0.25 0.25 NA 0.75 1.00  
cume_dist(y)
```

```
## [1] 0.2 0.6 0.6 NA 0.8 1.0
```

```
#> [1] 0.2 0.6 0.6 NA 0.8 1.0
```

## summarise()

```
summarise(flights, delay = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 1 x 1  
##   delay  
##   <dbl>  
## 1  12.6
```

```
by_day <- group_by(flights, year, month, day)
summarise(by_day, delay = mean(dep_delay, na.rm = TRUE))
```

```
## `summarise()` regrouping output by 'year', 'month' (override with `.groups` argument)
```

```
## # A tibble: 365 x 4
## # Groups:   year, month [12]
##   year month   day delay
##   <int> <int> <int> <dbl>
## 1  2013     1     1  11.5
## 2  2013     1     2  13.9
## 3  2013     1     3  11.0
## 4  2013     1     4   8.95
## 5  2013     1     5   5.73
## 6  2013     1     6   7.15
## 7  2013     1     7   5.42
## 8  2013     1     8   2.55
## 9  2013     1     9   2.28
## 10 2013     1    10   2.84
## # ... with 355 more rows
```


**pipe** %>%

```
delays <- flights %>%
  group_by(dest) %>%
  summarise(
    count = n(),
    dist = mean(distance, na.rm = TRUE),
    delay = mean(arr_delay, na.rm = TRUE)
  ) %>%
  filter(count > 20, dest != "HNL")
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
ggplot(data = delays, mapping = aes(x = dist, y = delay)) +
  geom_point(aes(size = count), alpha = 1/3) +
  geom_smooth(se = FALSE)
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
flights %>%  
  group_by(year, month, day) %>%  
  summarise(mean = mean(dep_delay))
```

```
## `summarise()` regrouping output by 'year', 'month' (override with `.groups` argumen  
t)
```

```
## # A tibble: 365 x 4
## # Groups:   year, month [12]
##   year month   day mean
##   <int> <int> <int> <dbl>
## 1  2013     1     1    NA
## 2  2013     1     2    NA
## 3  2013     1     3    NA
## 4  2013     1     4    NA
## 5  2013     1     5    NA
## 6  2013     1     6    NA
## 7  2013     1     7    NA
## 8  2013     1     8    NA
## 9  2013     1     9    NA
## 10 2013     1    10    NA
## # ... with 355 more rows
```

without `na.rm`, we will get a lot of missing values. (If one of the input is NA, the calculation result will be NA.)

**Useful summary functions** \* `mean(x)`, `median(x)`,

- `sd(x)` standard deviation, `IQR(x)` interquartile range, `mad(x)` median absolute deviation,
- `min(x)`, `quantile(x, 0.25)`, `max(x)`
- `first(x)`, `nth(x,2)`, `last(x)`
- `n()` count, `n_distinct(x)`, `sum(!is.na(x))`
- `sum(x < 10)`, `mean(y==0)`, converts TRUE to 1 and FALSE to 0,

remember to **ungrouping** `ungroup()`